KESTRELATTICE

Governed Agent Mesh | Micro-Product

The Playbook -- Studio Edition (Free Sample)

Free sample edition of The Playbook -- Studio Edition by ghostcorpnet. Sections 1-2 in full: why governance is the hard part, and the 8-layer governance stack for AI agents. Full edition: https://koalstin.gumroad.com/l/hdigmr ($29, no refunds).

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

The Playbook -- Studio Edition

Free Sample Edition

Running multi-agent AI operations with identity, spend limits, approval gates, and audit trails.

By GhostCorp · ghostcorpnet · 2026

This free sample contains the complete first two sections of The Playbook -- Studio Edition. The full edition ($29) adds the implementation patterns, tier tables, spend-limit configurations, audit-trail schemas, and the two-week rollout plan.

Full edition: https://koalstin.gumroad.com/l/hdigmr

Please note: ghostcorpnet sells digital products with a no-refunds policy, stated plainly on every listing.

1. Why governance, not orchestration, is the hard part

Orchestration is a solved problem: LangGraph, CrewAI, AutoGen, OpenAI Swarm and a dozen other frameworks will happily run a swarm of agents for you. What none of them give you is the answer to three questions your business depends on:

1. **What did the agents do?**
2. **What did they spend?**
3. **What were they allowed to do?**

Practitioners feel this daily. A developer on r/IndieDev documented an infinite-loop agent burning **$700+ in API credits in 72 hours** because nothing in the stack had budget controls or a watchdog. Threads on r/AI_Agents keep asking the same two things: why do frameworks let agents rack up unauthorized costs with no audit trail, and who gave my agent authority to do that? Builders on Hacker News argue that prompt engineering is inadequate and real systems need a **centralized safety mesh control plane**.

Meanwhile the market is not waiting. Gartner forecasts **40% of enterprise applications embedding task-specific AI agents by end-2026** (up from <5% in 2025) and **$206.5B in AI-agent software spend in 2026**. Enterprises are already paying for this class of control: LangSmith at **$39/seat/mo**, Langfuse up to **$2,499/mo**, AgentOps enterprise from **$2,000/mo**.

This playbook is the governance layer we built for ghostcorpnet, written down so you can apply it to any agent stack.

2. The 8-layer governance stack

| Layer | What it guarantees | Fail mode if missing
1 | Identity | Every agent action is attributable to an owner-scoped identity | Unverifiable actions, no accountability
2 | Grants | Explicit, least-privilege, time-bounded permission records | Ambient authority, scope creep
3 | Approval gates | Risk-tiered human decisions on money/outreach/contracts/permissions | Runaway external side effects
4 | Spend limits | Hard budget ceilings per plan, workload, and expiry window | The $700 infinite loop
5 | Fail-closed runtime | Missing evidence blocks the action instead of warning | Silent degradation into unsafe behavior
6 | Audit trails | Hash-linked rows for every request, approval, block, outcome | "Trust me, it worked"

7 | Telemetry | What ran, how long, at what cost, bounded in scope | Invisible drift and cost leaks
8 | Owner control | Secrets never stored in records; terms never auto-accepted | Leaked credentials, accidental contract acceptance

The order matters: each layer assumes the ones above it. You cannot have meaningful spend limits (4) without attributable identity (1), and your audit trail (6) is worthless if actions can bypass grants (2).

What the full edition adds

The complete Studio Edition continues with the working patterns behind every layer:

- **3. Identity and authority: the pattern** -- owner-scoped identities and grant records, and the rule that parks any capability request with no recorded identity as REQUIRES_OWNER_ACTION.
- **4. Risk-tiered approval gates** -- the three-tier policy table (hard-gated / bounded auto / draft-only) and why you must record the rejection path.
- **5. Spend limits that actually stop the bleeding** -- ceilings per plan, watchdog patterns, expiry as a spend limit, fail-closed budgets.
- **6. Audit trails: hash-linked and claim-evidenced** -- the audit-row schema and the discipline that makes trails trustworthy.
- **7. Fail-closed as a design philosophy** -- converting every warning path into a block with a recorded reason and an unblock checklist.
- **8. The $0 open-source distribution stack** -- the verified free-tier tooling plan, with honest funnel benchmarks.
- **9. Rollout plan: governance in two weeks** -- the day-by-day sequence (identity, grants + audit, spend limits, approval gates, fail-closed sweep).
- **10. Where the market is going** -- the funding and analyst signals behind the governance layer.

Get the full edition ($29): https://koalstin.gumroad.com/l/hdigmr

Also available: the Stop-the-Bleed Pair -- the cost-control workbook + risk audit kit bundled: https://koalstin.gumroad.com/l/stop-the-bleed-pair

All ghostcorpnet products are digital downloads with a no-refunds policy.

This sample is an excerpt of an original GhostCorp work. Every factual claim in the full edition links to a verifiable public source in the buyer-evidence ledger.

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com