KESTRELATTICE

Governed Agent Mesh | Micro-Product

Agent Bounds Checklist (Free)

Free one-page checklist from ghostcorpnet: the bounds every AI agent plan needs before it runs -- workload, data class, resources, time, location, egress, storage, cost, expiry -- plus approval tiers and stop conditions. From The Playbook -- Studio Edition: https://koalstin.gumroad.com/l/hdigmr

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

Agent Bounds Checklist

Before any AI agent runs, state its bounds. If you can't state the bounds, you don't have a plan -- you have a wish.

From The Playbook -- Studio Edition by ghostcorpnet. Free to use and share.

The nine bounds (fill in every one)

- [ ] **Workload** -- exactly what tasks this agent may run, and nothing else
- [ ] **Data class** -- what data it may read or write (public / internal / customer / secrets)
- [ ] **Resources** -- which tools, APIs, and systems it may touch
- [ ] **Time** -- when it may run, and the hard stop time
- [ ] **Location** -- where it runs and where its outputs may go
- [ ] **Egress** -- what it may send outside (email, posts, API calls) and to whom
- [ ] **Storage** -- what it may create, keep, or delete
- [ ] **Cost ceiling** -- the maximum spend for this plan, in dollars, counted per plan (not per account)
- [ ] **Expiry** -- the date/time this plan dies unless a human renews it

Identity and grants

- [ ] Every agent has a stable, owner-scoped identity -- no anonymous or shared actors
- [ ] Every capability is an explicit grant record: capability, scope, expiry
- [ ] Least privilege by default; a request with no recorded identity is blocked, not warned

Approval tiers

- [ ] **Hard-gated (always a human):** money, contracts, external outreach, credentials, permission changes, legal actions
- [ ] **Bounded auto:** routine work inside the nine bounds above -- runs automatically, fully audited, auto-blocks at any bound
- [ ] **Draft-only:** proposals and messages are prepared and recorded, never executed

Spend controls

- [ ] A meter or proxy counts spend in front of the framework -- something is always counting
- [ ] When the ceiling is hit, the result is **blocked and recorded as blocked** -- never retried, never billed

Stop conditions

- [ ] A kill condition is named in advance (error rate, spend rate, unexpected egress, owner command)
- [ ] A blocked action persists an exact blocked result -- "it probably didn't run" is not an audit trail
- [ ] Every action writes an audit row: actor, action, target, outcome, timestamp

Want the full system? The Playbook -- Studio Edition ($29) covers the 8-layer governance stack, approval-gate patterns, spend-limit configurations, and hash-linked audit trails, with a two-week rollout plan:

https://koalstin.gumroad.com/l/hdigmr

ghostcorpnet digital products carry a no-refunds policy, stated plainly on every listing.

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com